

Session ID #EGA77B878S

Foundational: Ansible / Terraform



Noel Colon
Garage Solution Engineering

John Webb
Garage Solution Engineering

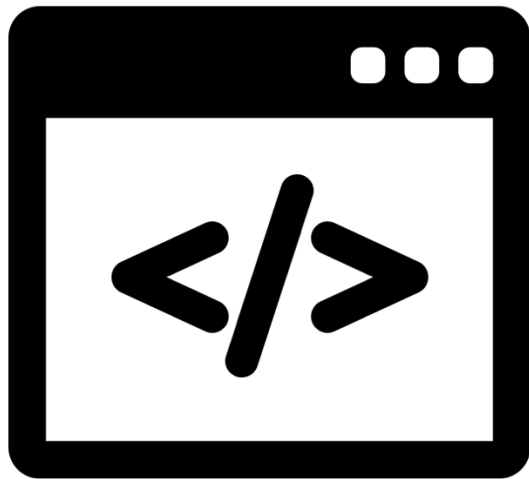


Legal Disclaimer

© IBM Corporation 2020. All Rights Reserved.

The information contained in this publication is provided for informational purposes only. While efforts were made to verify the completeness and accuracy of the information contained in this publication, it is provided AS IS without warranty of any kind, express or implied. In addition, this information is based on IBM's current product plans and strategy, which are subject to change by IBM without notice. IBM shall not be responsible for any damages arising out of the use of, or otherwise related to, this publication or any other materials. Nothing contained in this publication is intended to, nor shall have the effect of, creating any warranties or representations from IBM or its suppliers or licensors, or altering the terms and conditions of the applicable license agreement governing the use of IBM software.

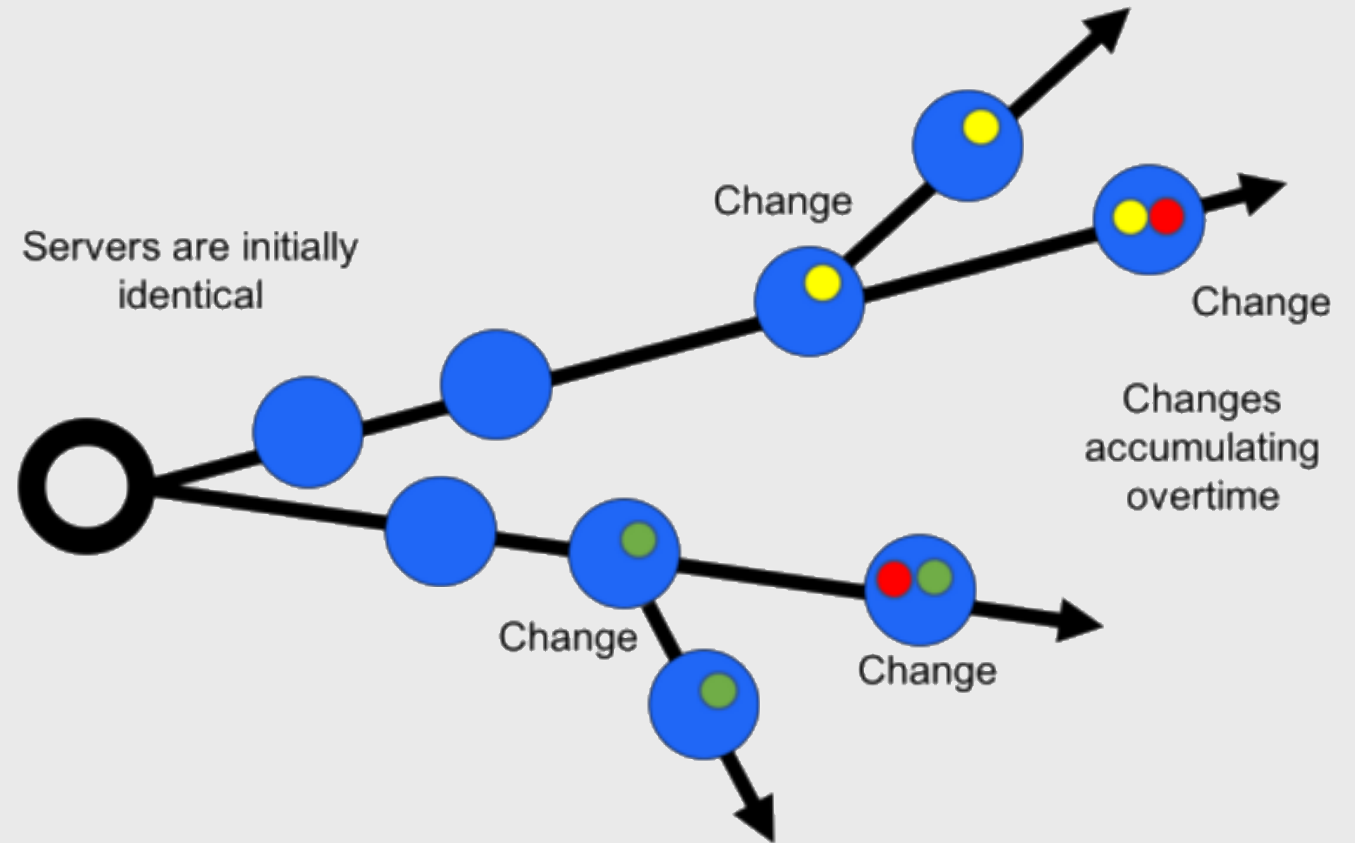
References in this presentation to IBM products, programs, or services do not imply that they will be available in all countries in which IBM operates. Product release dates and/or capabilities referenced in this presentation may change at any time at IBM's sole discretion based on market opportunities or other factors, and are not intended to be a commitment to future product or feature availability in any way. Nothing contained in these materials is intended to, nor shall have the effect of, stating or implying that any activities undertaken by you will result in any specific sales, revenue growth or other results.

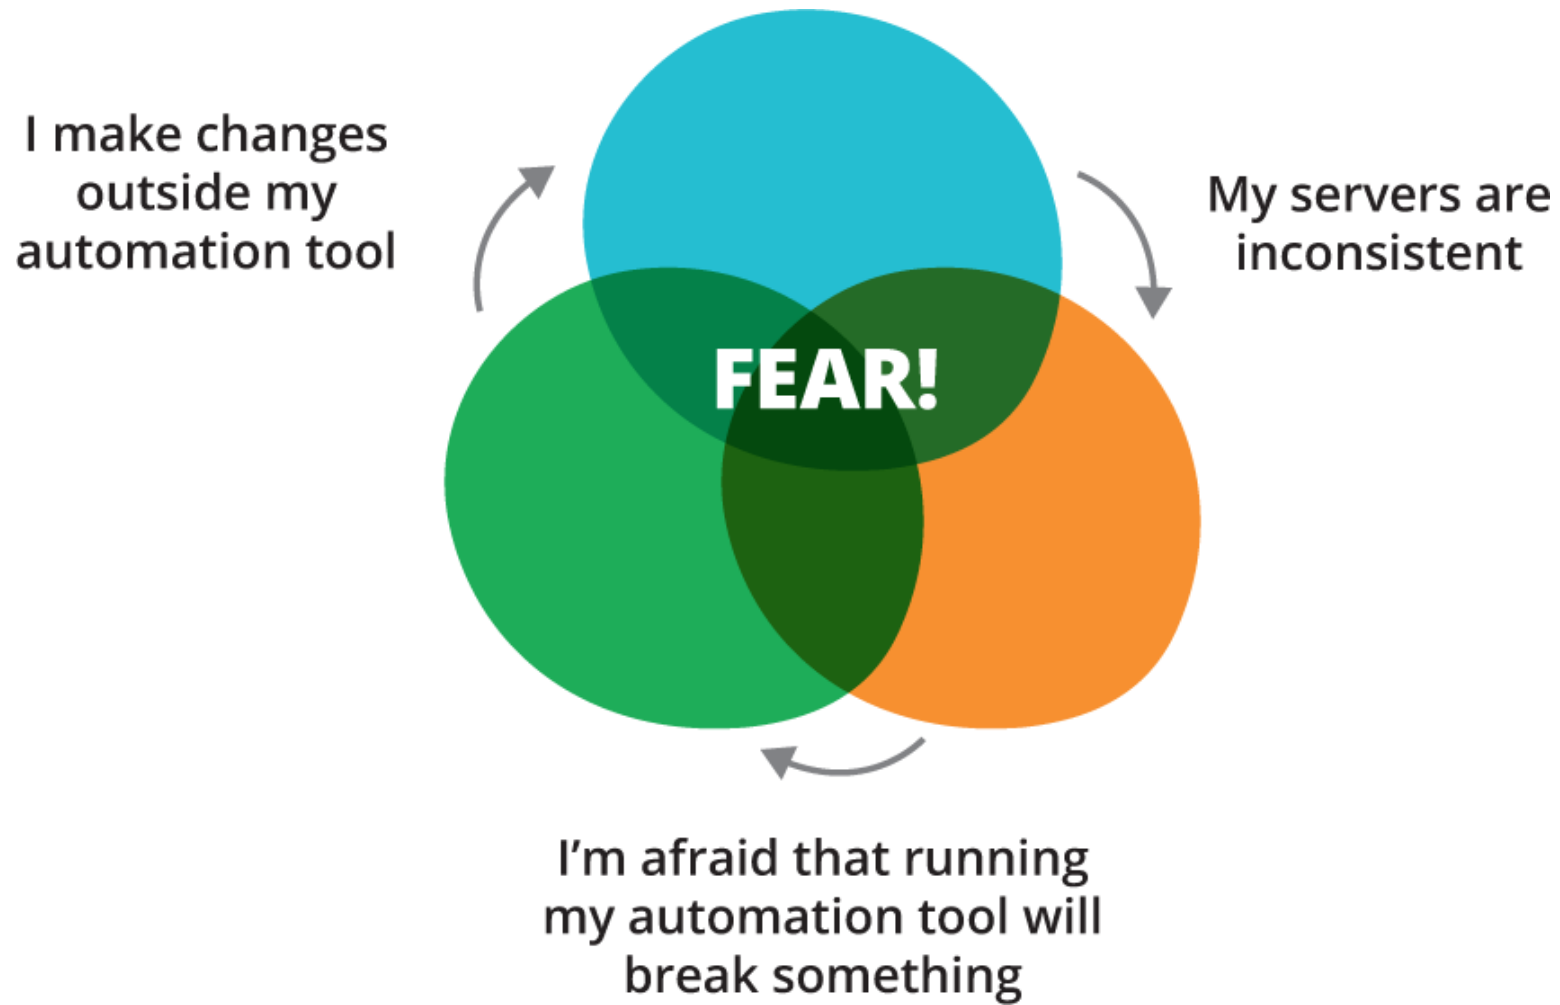


Infrastructure as code (IaC) is the process of managing and provisioning computer data centers through machine-readable definition files, rather than physical hardware configuration or interactive configuration tools.

The value of IaC can be broken down into three measurable categories: cost (reduction), speed (faster execution) and risk (remove errors and security violations).[1](#)

Configuration Drift





The Automation Fear Spiral

What's infrastructure ?



Traditionally the Cloud services, hosts, virtual machines, Docker containers, networking (*routing, switching, firewalls*), and storage were considered infrastructure. Now infrastructure also includes more complex services or Software-as-a-Service products delivered by third parties such as DNS, Content Delivery Networks (*CDN*), databases, job scheduling, queues, K8s, monitoring.... Terraform can configure components like GitHub organizations and repositories, Grafana monitoring console,.....

What's Configuration Management ?

Configuration management is a systems engineering process for establishing and maintaining consistency of a product's performance, functional, and physical attributes with its requirements, design, and operational information throughout its life.



**USING THE RIGHT TOOL
SAVES**

Make the right
choice

LOTS OF PAIN





HashiCorp

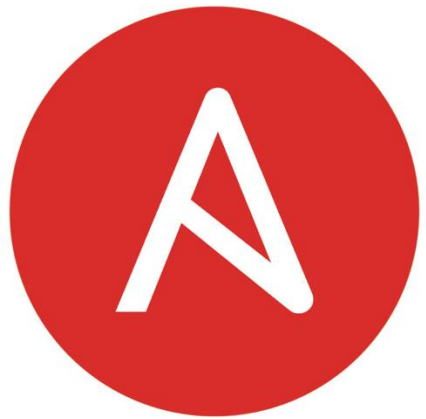
Terraform

Write, Plan, and Create Infrastructure as
Code

Terraform is a tool for building, changing, and versioning infrastructure safely and efficiently

Terraform generates an execution plan describing what it will do to reach the desired state, and then executes it to build the described infrastructure. As the configuration changes, Terraform is able to determine what changed and create incremental execution plans which can be applied.

- <https://www.terraform.io/>
- <https://github.com/hashicorp/terraform>



RED HAT[®]
ANSIBLE[®]
Automation

Ansible is an open-source software provisioning, configuration management, and application-deployment tool

It uses a declarative configuration language: you declare your desired state, and Ansible will manage how to get your systems to that state

➤ <https://www.ansible.com/>

Terraform Primer

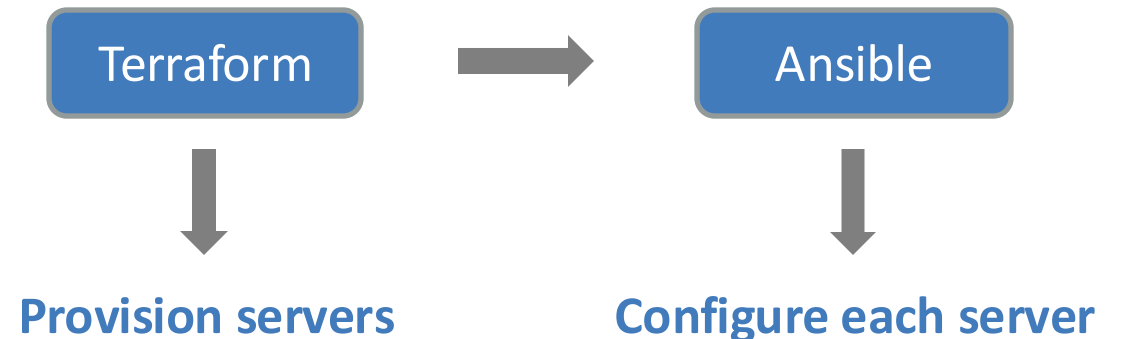
Providers, Data Sources and Resources

Declarative definitions style to define resources to provision

The Terraform and Ansible configuration language is declarative, describing an intended goal rather than the steps to reach that goal.

- ✓ With declarative definitions you specify what should be there and not how to do
- ✓ Only missing parts must be created, existing ones must be in the desired state, and obsolete ones must be destroyed
- ✓ Ensure the correct order of creation
- ✓ Integrate with other tools such as Ansible
- ✓ Platform agnostic
- ✓ Update management
- ✓ Extension capabilities

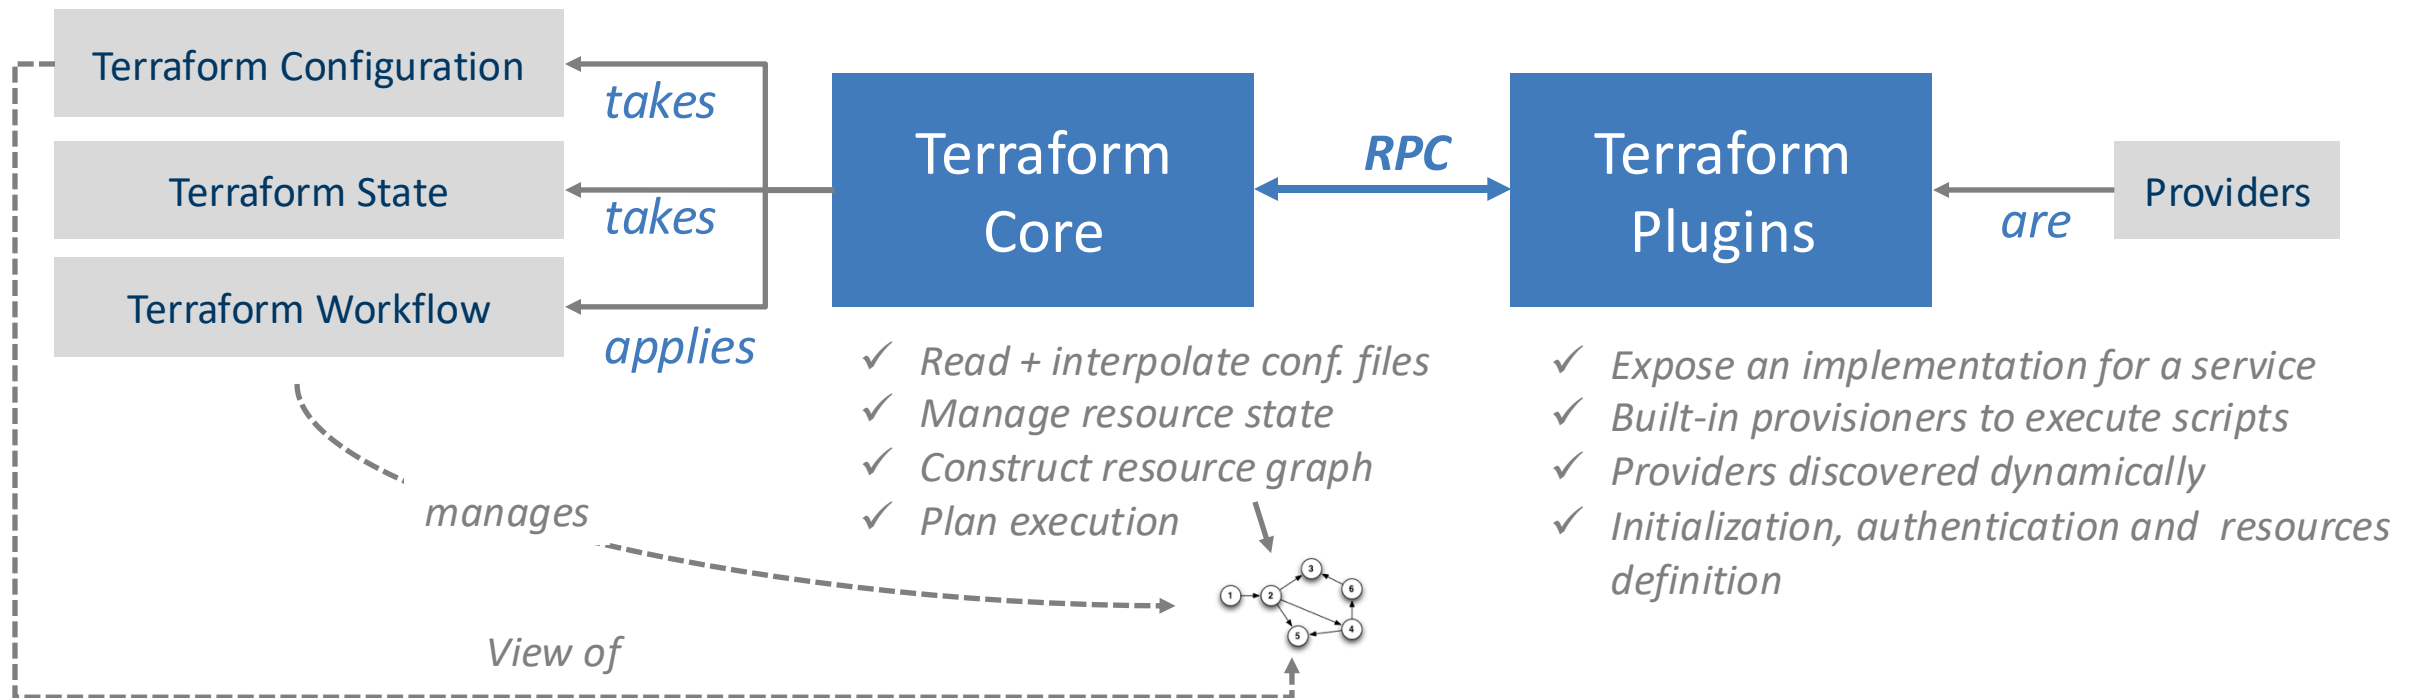
Declarative style



How Terraform operates ?

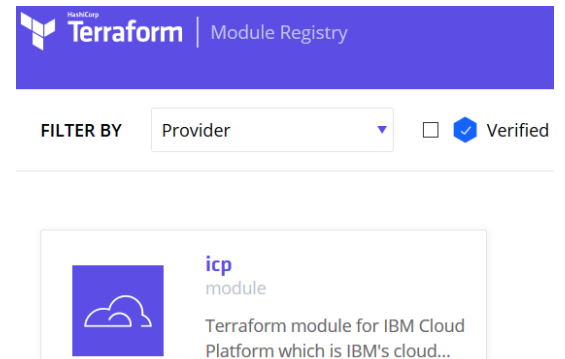
Terraform examines each resource and uses a graph-based approach to model and apply the desired state.

Each resource is placed inside the graph, its relationships with other resources are calculated, and then each resource is automatically built in the correct order to produce your infrastructure.



Introduction & Concepts

- You define **resources** as code in Terraform templates
 - ✓ Specify the provider
 - ✓ Specify provisioners
 - ✓ Specify the resources
 - ✓ Parameterize your template using **variables**
- **Provider**: a source of resources with API endpoint & authentication
- **Provisioner** execute local or remote scripts during resource creation or destroy time
- **Data Source**: information read from provider
- **Modules** allows to reuse the same code in different environments
 - ✓ Modules should be sourced from git tags / branches
 - ✓ Module registry
- You follow a **Workflow**:
 - ✓ **Plan** => see what you are about to deploy
 - ✓ **Apply** => to apply the changes
- Terraform generates a **state file**:
 - ✓ Store information about your managed infrastructure and configuration (*apply execution result*)
 - ✓ Generated during the apply stage, don't edit manually
- A **backend** determines how a state is loaded and how operations are executed



Terraform Syntax

Argument

Assign a value to a name
`image_id = "abc123"`

Block

Container for other content

has

Block Type

Block Body

{ }

type

labels

```
resource "aws_instance" "example" {  
  ami = "abc123"  
  
  network_interface {  
    # ...  
  }  
}
```

Top level block type

- ✓ Resource
- ✓ Input variable
- ✓ Output value
- ✓ Data source

comment

```
variable "variable_name" {  
  type = "variable_type"  
  default = "variable_default_value" # optional  
}
```

- ✓ string
- ✓ map
- ✓ list
- ✓ boolean



```
resource "aws_instance" "web" {  
  ami = "${var.ami}"  
  instance_type = "t2.micro"  
  
  tags {  
    Name = "HelloWorld"  
  }  
}
```

Variable interpolation

- ✓ Evaluate the expression
- ✓ Convert the value to a string
- ✓ Insert it into the string

Providers

<https://www.terraform.io/docs/configuration/providers.html>

```
provider "vsphere" {  
  version = "~> 1.11.0"  
  vsphere_server = "${var.vsphere_server}"  
  
  # if you have a self-signed cert  
  allow_unverified_ssl = "${var.allow_unverified_ssl}"  
}
```

- The provider block is used to configure the named provider
- A provider is responsible for creating and managing resources
- Multiple provider blocks can exist if a Terraform configuration is composed of multiple providers
 - Credential can be passed in different ways (provider dependent)

Data Sources

Used to discover
information about
existing objects

Can be nested in on other
data objects

```
data "vsphere_datacenter" "dc" {
    name = "${var.vsphere_datacenter}"
}

data "vsphere_datastore_cluster" "datastore_cluster" {
    name = "${var.datastore_cluster}"
    datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "vsphere_resource_pool" "pool" {
    name = "${var.vsphere_cluster}/Resources/${var.vsphere_resource_pool}"
    datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "vsphere_network" "network" {
    name = "${var.network_label}"
    datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "vsphere_virtual_machine" "template" {
    name = "${var.template}"
    datacenter_id = "${data.vsphere_datacenter.dc.id}"
}
```

Resources

Defines a resource that exists within the infrastructure

May be a physical component such as an EC2 instance, or it can be a logical resource such as an SSH key

Resource definition is described by the provider

```
resource "vsphere_virtual_machine" "camlab" {
  name = "terraform-test"
  resource_pool_id = "${data.vsphere_resource_pool.pool.id}"
  datastore_id = "${data.vsphere_datastore.datastore.id}"

  num_cpus = 2
  memory = 1024
  guest_id = "other3xLinux64Guest"

  network_interface {
    network_id = "${data.vsphere_network.network.id}"
  }

  disk {
    label = "disk0"
    size = 20
  }
}
```

Defining Variables

All input variables must be defined in variable blocks

Multiple types of variables string, map, list, etc.

Can define a default or leave default blank

Variables that are not defined will need to be defined a run time

```
variable "vsphere_server" {  
    description = "vsphere server to connect to"  
    default = "10.0.0.210"  
}
```

```
variable "camlab" {  
    type = "map"  
  
    default = {  
        nodes = "1"  
        vcpu = "2"  
        memory = "4096"  
    }  
}
```

```
variable "dns_servers" {  
    description = "DNS Servers to configure on VMs"  
    default = ["8.8.8.8", "8.8.4.4"]  
}
```

Referencing Variables

Input variables are referenced in the “\${var.<name>}” syntax

Data variables are referenced in the “\${data.<name>}” syntax

Local variables are referenced in the “\${local.<name>}” syntax

```
vsphere_server = "${var.vsphere_server}"
```

```
num_cpus = "${var.camlab["vcpu"]}"
```

```
path = "${local.team_folder}"
```

Terraform Primer

Project layout

Template planning

terraform plan

- Connects to the hypervisor and performs a dry run to show you what will be created if you apply the template
- Will perform syntax checks and error if you have errors in your code

```
export VSPHERE_USER=administrator@vsphere.local  
export VSPHERE_PASSWORD=Passw0rd!
```

Note: You will need to provide credentials to connect your hypervisor

```
Refreshing Terraform state in-memory prior to plan...  
The refreshed state will be used to calculate this plan, but will not be  
persisted to local or remote state storage.
```

```
data.vsphere_datacenter.dc: Refreshing state...  
data.vsphere_resource_pool.pool: Refreshing state...  
data.vsphere_network.network: Refreshing state...  
data.vsphere_datastore_cluster.datastore_cluster: Refreshing state...  
data.vsphere_virtual_machine.template: Refreshing state...
```

```
-----  
An execution plan has been generated and is shown below.  
Resource actions are indicated with the following symbols:
```

```
+ create
```

```
Terraform will perform the following actions:
```

```
# tls_private_key.ssh will be created  
+ resource "tls_private_key" "ssh" {  
  + algorithm      = "RSA"  
  + ecdsa_curve    = "P224"  
  + id             = (known after apply)  
  + private_key_pem = (known after apply)  
  + public_key_fingerprint_md5 = (known after apply)  
  + public_key_openssh = (known after apply)  
  + public_key_pem   = (known after apply)  
  + rsa_bits         = 2048  
}
```

```
# vsphere_folder.icpenv[0] will be created  
+ resource "vsphere_folder" "icpenv" {  
  + datacenter_id = "datacenter-21"  
  + id            = (known after apply)  
  + path          = "Target/Team01/Lab2"  
  + type          = "vm"  
}
```

Applying your template

terraform apply

- Will first connect to the hypervisor and perform a plan
- Then prompts you to apply the changes that are displayed
- If approved it will apply the changes and create the resources defined

```
Do you want to perform these actions?  
Terraform will perform the actions described above.  
Only 'yes' will be accepted to approve.  
  
Enter a value: yes  
  
vsphere_virtual_machine.camlab[0]: Destroying... [id=42242bf6-329b-0c9d-234a-6ee84276353b]  
vsphere_virtual_machine.camlab[0]: Still destroying... [id=42242bf6-329b-0c9d-234a-6ee84276353b, 10s elapsed]  
vsphere_virtual_machine.camlab[0]: Destruction complete after 19s  
vsphere_virtual_machine.camlab[0]: Creating...  
vsphere_virtual_machine.camlab[0]: Still creating... [10s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [20s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [30s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [40s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [50s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [1m0s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [1m10s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [1m20s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [1m30s elapsed]  
vsphere_virtual_machine.camlab[0]: Still creating... [1m40s elapsed]  
vsphere_virtual_machine.camlab[0]: Creation complete after 1m50s [id=4224ee76-e1fd-25bd-fa82-5a39321be001]  
  
Apply complete! Resources: 1 added, 0 changed, 1 destroyed.
```

Plan: 3 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?

Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes

Terraform state files

tfstate

- Tracks the changes that have been made by terraform
- Used to define the declared state and will be compared to actual provider resource to determine configuration change
- JSON format can be queried using **terraform show**

```
drwxr-xr-x  3 root root   163 Jul  7 18:35 .
drwx----- 26 root root  4096 Jul  7 18:33 ..
-rw-r--r--  1 root root  4156 Jul  7 18:33 instances.tf
-rw-r--r--  1 root root   196 Jul  7 18:15 outputs.tf
drwxr-xr-x  3 root root    21 Jul  7 18:15 .terraform
-rw-r--r--  1 root root 12857 Jul  7 18:35 terraform.tfstate
-rw-r--r--  1 root root 12923 Jul  7 18:33 terraform.tfstate.backup
-rw-r--r--  1 root root  1184 Jul  7 18:15 terraform.tfvars
-rw-r--r--  1 root root  4301 Jul  7 18:17 variables.tf
```

```
{
  "version": 4,
  "terraform_version": "0.12.0",
  "serial": 6,
  "lineage": "043e9eea-b239-fb5f-c1d7-6def604feefa",
  "outputs": {},
  "resources": [
    {
      "mode": "data",
      "type": "vsphere_datacenter",
      "name": "dc",
      "provider": "provider.vsphere",
      "instances": [
        {
          "schema_version": 0,
          "attributes": {
            "id": "datacenter-21",
            "name": "Datacenter"
          }
        }
      ]
    },
    {
      "mode": "data",
      "type": "vsphere_datastore_cluster",
      "name": "datastore_cluster",
      "provider": "provider.vsphere",
      "instances": [
        {
          "schema_version": 0,
          "attributes": {
            "datacenter_id": "datacenter-21",
            "id": "group-p163",
            "name": "DatastoreCluster"
          },
          "depends_on": [
            "data.vsphere_datacenter.dc"
          ]
        }
      ]
    }
  ]
}
```

Template Deletion

terraform destroy

- Will first connect to the hypervisor compare what is there to the .tfstate file
- It will show you what is going to be deleted
- Then prompts you to destroy the resources that are displayed
- If approved it will destroy all the resources

Plan: 0 to add, 0 to change, 3 to destroy.

Do you really want to destroy all resources?

Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

```
# vsphere_folder.icpenv[0] will be destroyed
- resource "vsphere_folder" "icpenv" {
  - datacenter_id = "datacenter-21" -> null
  - id            = "group-v466"    -> null
  - path         = "Target/Team10/Lab2" -> null
  - type         = "vm"            -> null
}

# vsphere_virtual_machine.camlab[0] will be destroyed
- resource "vsphere_virtual_machine" "camlab" {
  - boot_delay                = 0 -> null
  - boot_retry_delay          = 10000 -> null
  - boot_retry_enabled        = false -> null
  - change_version            = "2019-07-08T01:41:36.94448Z" -> null
  - cpu_hot_add_enabled        = false -> null
  - cpu_hot_remove_enabled    = false -> null
  - cpu_limit                  = -1 -> null
  - cpu_performance_counters_enabled = false -> null
  - cpu_reservation           = 0 -> null
  - cpu_share_count            = 2000 -> null
  - cpu_share_level            = "normal" -> null
  - datastore_cluster_id       = "group-p163" -> null
  - datastore_id               = "datastore-107" -> null
  - default_ip_address         = "10.0.0.190" -> null
  - efi_secure_boot_enabled    = false -> null
  - enable_disk_uuid           = false -> null
  - enable_logging             = false -> null
  - ept_rvi_mode               = "automatic" -> null
  - firmware                   = "bios" -> null
  - folder                     = "Target/Team10/Lab2" -> null
  - force_power_off            = true -> null
  - guest_id                   = "rhel7_64Guest" -> null
  - guest_ip_addresses         = [
    - "10.0.0.190",
    - "fe80::250:56ff:fea4:4792",
  ] -> null
  - host_system_id             = "host-44" -> null
  - hv_mode                    = "hvAuto" -> null
  - id                         = "4224ee76-e1fd-25bd-fa82-5a39321be001" -> null
}
```

Advanced Topics

Interpolation, modules and providers

Modules

- A *module* is a container for multiple resources that are used together.
- Every Terraform configuration has at least one module, known as its *root module*, which consists of the resources defined in the .tf files in the main working directory.
- A module can call other modules, which lets you include the child module's resources into the configuration in a concise way. Modules can also be called multiple times, either within the same configuration or in separate configurations, allowing resource configurations to be packaged and re-used.

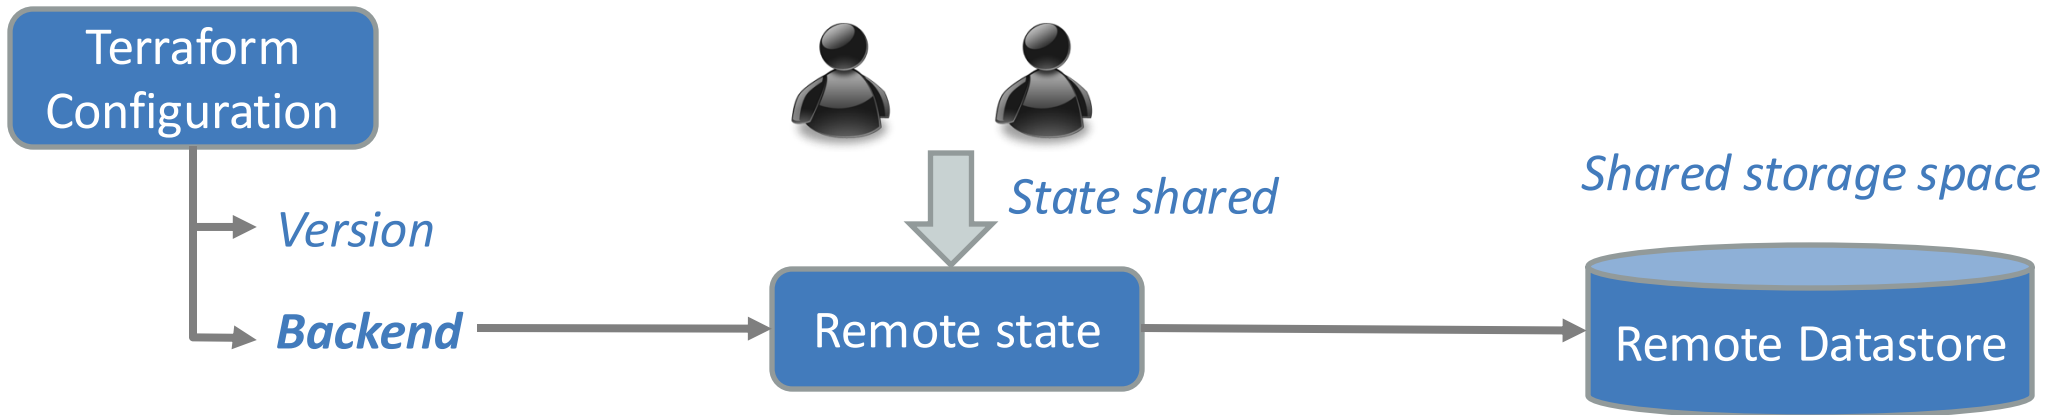
```
#####
### Deploy ICP to cluster
#####
module "icpprovision" {
  source = "github.com/ibm-cloud-architecture/terraform-module-icp-deploy.git?ref=2.3.6"

  # Provide IP addresses for master, proxy and workers
  boot-node = "${vsphere_virtual_machine.icpmaster.0.default_ip_address}"
  icp-host-groups = {
    master = ["${vsphere_virtual_machine.icpmaster.*.default_ip_address}"]
    proxy = ["${vsphere_virtual_machine.icpproxy.*.default_ip_address}"]
    worker = ["${vsphere_virtual_machine.icpworker.*.default_ip_address}"]
    // make the master nodes managements nodes if we don't have any specified
    management = "${slice(concat(vsphere_virtual_machine.icpmanagement.*.default_ip_address,
                                vsphere_virtual_machine.icpmaster.*.default_ip_address),
                                0, var.management["nodes"] > 0 ? length(vsphere_virtual_machine.icpmanagement.*.default_ip_address) : 0)}"
    va = ["${vsphere_virtual_machine.icpva.*.default_ip_address}"]
  }
}
```

<https://www.terraform.io/docs/configuration/modules.html>

Store your state file remotely using a terraform backend

- Protect state with locks to prevent corruption
- Init your backend: *terraform init*



- ✓ *Work in a team*
- ✓ *Team work => state locking*
- ✓ *Remote operations*
- ✓ *Keep sensitive information off disk*
- ✓ *Persistent data stored in the backend belongs to a workspace*

```
terraform {  
  backend "consul" {  
    address = "demo.consul.io"  
    scheme  = "https"  
    path    = "example_app/terraform_state"  
  }  
}
```

Provisioners are used to execute scripts on a local or remote machine as part of resource creation or destruction. Provisioners can be used to bootstrap a resource, cleanup before destroy, run configuration management, etc.

```
# Specify the ssh connection
connection {
    user = "${var.ssh_user}"
    password = "${var.ssh_password}"
    host = "${var.staticipblock != "0.0.0.0/0" ? cidrhost(var.staticipblock, (var.team_number * 10) + var.staticipblock_offset + count.index) : ""}"
}

provisioner "file" {
    source = "${path.module}/scripts"
    destination = "/tmp/terraform_scripts"
}

provisioner "remote-exec" {
    inline = [
        "sudo chmod u+x /tmp/terraform_scripts/*.sh",
        "/tmp/terraform_scripts/add-public-ssh-key.sh \"${tls_private_key.ssh.public_key_openssh}\"",
        "/tmp/terraform_scripts/add-private-ssh-key.sh \"${tls_private_key.ssh.private_key_pem}\" \"${var.ssh_user}\""
    ]
}
```

Provisioners

<https://www.terraform.io/docs/provisioners/index.html>

Ansible Primer

Inventory, Variables, Playbooks, Tasks

Inventory

Used to declare and
group hosts together

Default ansible inventory
/etc/ansible/hosts

```
mail.example.com
```

```
[webservers]  
foo.example.com  
bar.example.com  
web[01:10].example.com
```

```
[dbservers]  
one.example.com  
two.example.com  
three.example.com  
db-[a:f].example.com
```

Groups

Collection of hosts that share similar functions

mail.example.com

[web]

web01.example.com

web02.example.com

[db]

db01.example.com

db02.example.com

[dev]

web01.example.com

db01.example.com

[prod]

web02.example.com

db02.example.com

Special Groups:

all: every host in inventory

unbound: any host not in a group

Variables

variable values that relate
to a specific host or group
in inventory

```
mail.example.com
```

```
[web]
```

```
web01.example.com
```

```
web02.example.com
```

```
[db]
```

```
db01.example.com
```

```
db02.example.com
```

```
[web:vars]
```

```
web_root=/var/www/html
```

```
[db:vars]
```

```
db_name=database_name
```

Playbooks

The basis for a really simple configuration management and multi-machine deployment system, and one that is very well suited to deploying complex applications.

They declare configurations, but they can also orchestrate steps of any manual ordered process, even as different steps must bounce back and forth between sets of machines in particular orders.

```
---  
- hosts: web  
  remote_user: virtuser  
  become: yes  
  tasks:  
    - name: Gather the rpm package facts  
      package_facts:  
        manager: auto  
    - name: ensure apache is at the latest version  
      yum:  
        name: httpd  
        state: latest  
        when: "'httpd' not in ansible_facts.packages"  
    - name: make sure apache is running  
      service:  
        name: httpd  
        state: started
```

Executing Playbooks

```
ansible-playbook -i /path/to/inventory /path/to/playbook.yaml
```

- You can specify a specific inventory with the -i parameter

```
PLAY [web] *****
```

```
TASK [Gathering Facts] *****
```

```
ok: [web01.example.com]
```

```
ok: [web02.example.com]
```

```
TASK [Gather the rpm package facts] *****
```

```
ok: [web01.example.com]
```

```
ok: [web02.example.com]
```

```
TASK [ensure apache is at the latest version] *****
```

```
changed: [web01.example.com]
```

```
changed: [web02.example.com]
```

```
TASK [make sure apache is running] *****
```

```
changed: [web01.example.com]
```

```
changed: [web02.example.com]
```

```
PLAY RECAP *****
```

```
web01.example.com      : ok=4    changed=2    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
```

```
web02.example.com      : ok=4    changed=2    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
```

State Management

```
ansible-playbook -i /path/to/inventory /path/to/playbook.yaml
```

- What happens if I execute the same playbook again?

```
PLAY [web] *****
```

```
TASK [Gathering Facts] *****
```

```
ok: [web01.example.com]
```

```
ok: [web02.example.com]
```

```
TASK [Gather the rpm package facts] *****
```

```
ok: [web01.example.com]
```

```
ok: [web02.example.com]
```

```
TASK [ensure apache is at the latest version] *****
```

```
skipping: [web01.example.com]
```

```
skipping: [web02.example.com]
```

```
TASK [make sure apache is running] *****
```

```
changed: [web01.example.com]
```

```
ok: [web02.example.com]
```

```
PLAY RECAP *****
```

```
web01.example.com      : ok=3    changed=1    unreachable=0    failed=0    skipped=1    rescued=0    ignored=0
```

```
web02.example.com      : ok=3    changed=0    unreachable=0    failed=0    skipped=1    rescued=0    ignored=0
```

Ansible + Terraform

```
module.runplaybook.null_resource.run_playbook_create[0]: Creating...
module.runplaybook.null_resource.run_playbook_create[0]: Provisioning with 'remote-exec'...
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): Connecting to remote host via SSH...
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec):   Host: 9.42.67.220
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec):   User: virtuser
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec):   Password: true
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec):   Private key: true
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec):   Certificate: false
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec):   SSH Agent: true
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec):   Checking Host Key: false
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): Connected!
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ++ sudo grep requiretty /etc/sudoers
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ++ echo 1
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): + export ANSIBLE_SSL_PIPELINING=1
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): + ANSIBLE_SSL_PIPELINING=1
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): + /tmp/ansible_chroot.sh ansible-playbook -f 20 -i /tmp/playbook_51029816/ansible.cfg /tmp/playbook_51029816/playbooks/playbook.yaml

module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): PLAY [all] *****

module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): TASK [Gathering Facts] *****
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ok: [9.42.67.221]
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ok: [9.42.67.220]

module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): TASK [Gather the rpm package facts] *****
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ok: [9.42.67.221]
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ok: [9.42.67.220]

module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): TASK [ensure apache is at the latest version] *****
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): skipping: [9.42.67.220]
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): skipping: [9.42.67.221]

module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): TASK [make sure apache is running] *****
module.runplaybook.null_resource.run_playbook_create[0]: Still creating... [10s elapsed]
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ok: [9.42.67.221]
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): ok: [9.42.67.220]

module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): PLAY RECAP *****
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): 9.42.67.220      : ok=3    changed=0    unreachable=0    failed=0    skipped=1    rescued=0    ignored=0
module.runplaybook.null_resource.run_playbook_create[0] (remote-exec): 9.42.67.221      : ok=3    changed=0    unreachable=0    failed=0    skipped=1    rescued=0    ignored=0
```

<https://github.com/ibm-cloud-architecture/terraform-ansible-runplaybooks>

Code walkthrough

Simple to Complex

Code Samples

Basic Example

<https://github.ibm.com/john-webb/cam4admins/tree/master/Terraform/Lab2>

Advanced Examples (OCP in different Cloud providers)

<https://github.com/ibm-cloud-architecture/terraform-openshift4-aws>

<https://github.com/ibm-cloud-architecture/terraform-openshift4-azure>

<https://github.com/ibm-cloud-architecture/terraform-openshift4-gcp>



Enterprise Environments

IBM Schematics

Cloud Automation Manager

Ansible Tower

IBM Schematics

IBM Cloud Schematics delivers Terraform-as-a-Service so that you can use a high-level scripting language to model the resources that you want in your IBM Cloud environment, and enable Infrastructure as Code

Workspace metadata

Workspace name

terraform-openshift4-aws

Tags

provider:aws

Resource group

None

Description (optional)

Import your Terraform template

Don't have an existing Terraform template? Try out one of our [sample templates](#).

GitHub or GitLab repository URL

URL of the GitHub or GitLab repository that hosts your Terraform configuration files. This can be the root of the repository, or it can include a branch and folder.

https://github.com/ibm-cloud-architecture/terraform-openshift4-aws

Personal access token (private repositories only)

Personal access token to authenticate with private GitHub or GitLab repositories

Retrieve input variables

Retrieve input variables

Input variables

Name	Description	Type	Default	Override value	Sensitive
aws_region		string	us-east-2		<input type="checkbox"/>
aws_azs		list	["a", "b", "c"]		<input type="checkbox"/>
default_tags		map	{ }		<input type="checkbox"/>
infrastructure_id		string			<input type="checkbox"/>
clustername		string			<input type="checkbox"/>

Generate plan

Apply plan

Recent activity

Workspace created

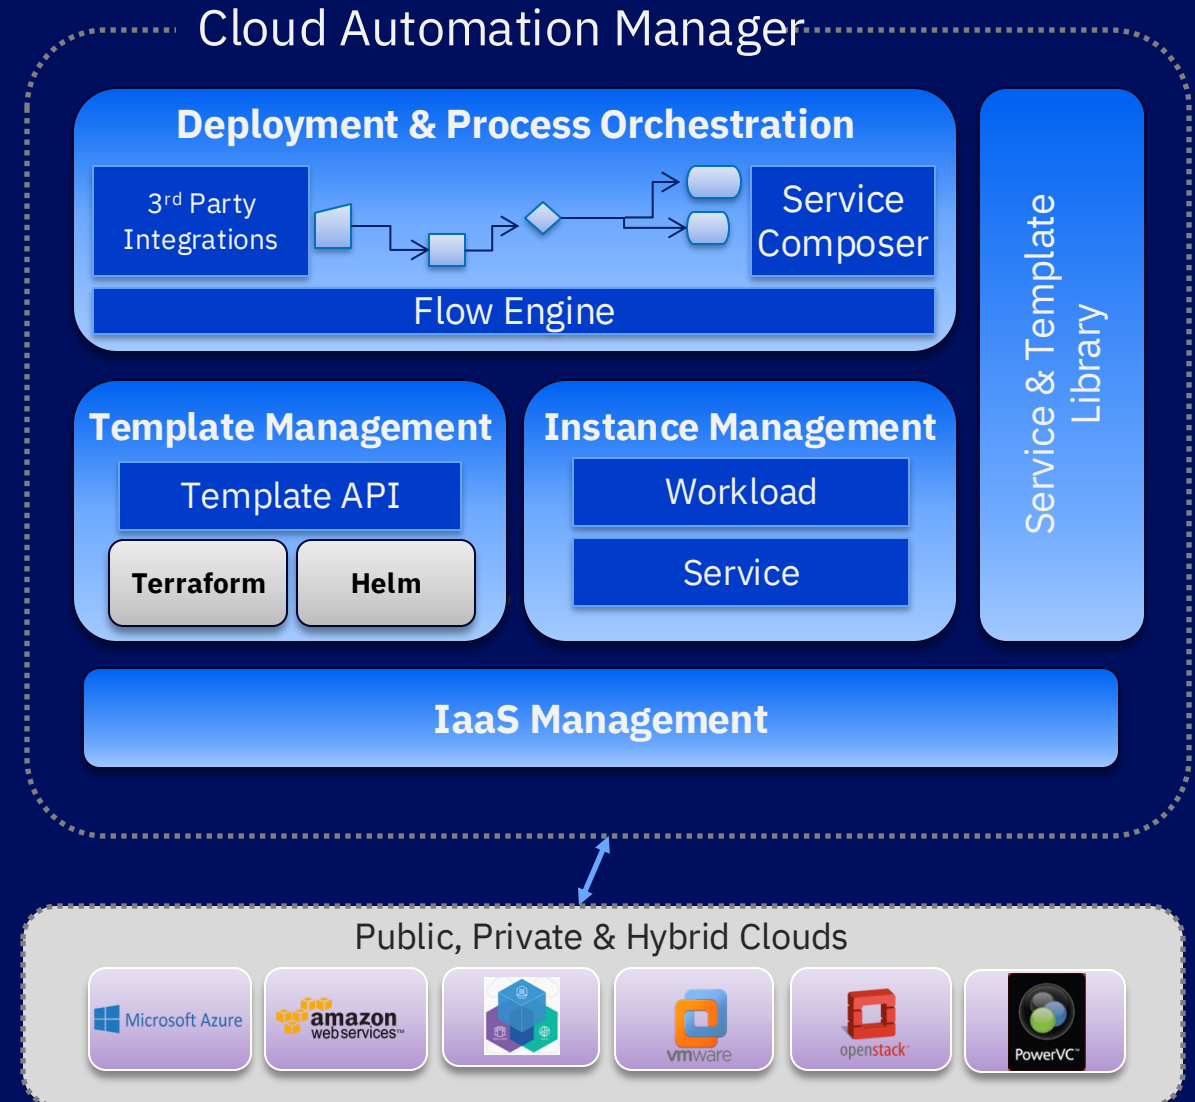
1/28/2020, 9:43:20 PM

<https://cloud.ibm.com/docs/schematics>

IBM Cloud Automation Manager

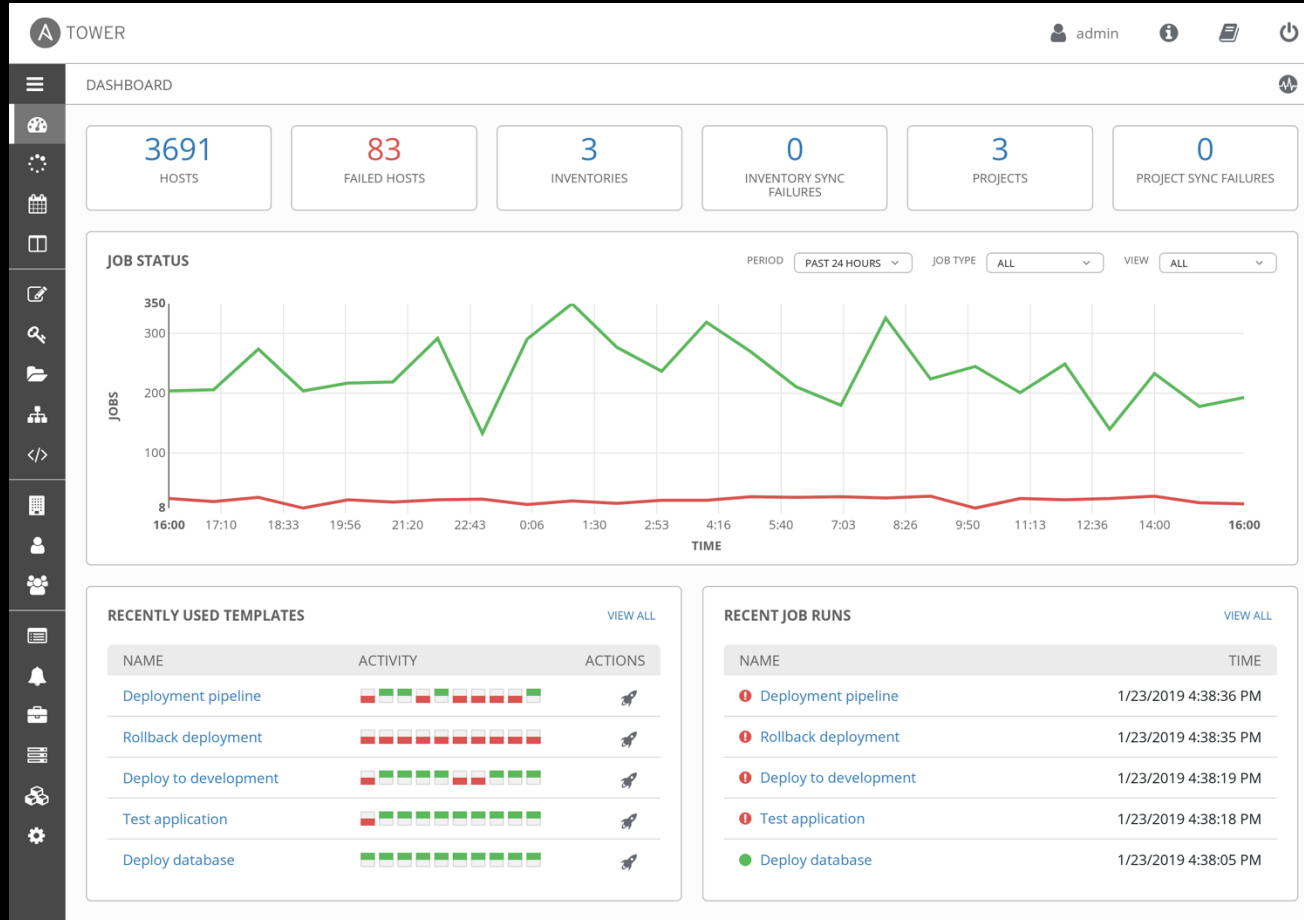
Full stack automation and service orchestration

- **Automated provisioning** – Automated provisioning of infrastructure and applications with workflow orchestration
- **Self-service** - Self-service access to cloud infrastructure and application services
- **Manage and govern** – Manage and govern workloads across multiple and hybrid clouds
- **Built with open technology** to avoid vendor lock-in



RedHat Ansible Tower

With Red Hat® Ansible® Tower you can centralize and control your IT infrastructure with a visual dashboard, role-based access control, job scheduling, integrated notifications and graphical inventory management



Dashboard

NOC-style display with access to all your environment data, recent activity, etc

<https://www.ansible.com/products/tower>

RedHat Ansible Tower

The screenshot displays the RedHat Ansible Tower web interface. The top navigation bar includes the 'TOWER' logo, a user profile 'admin', and icons for help, settings, and power. The main content area is titled 'JOBS / 29170 - Deploy application'. On the left, a 'DETAILS' sidebar shows job information: STATUS (Successful), STARTED (1/23/2019 6:52:19 PM), FINISHED (1/23/2019 6:52:35 PM), JOB TEMPLATE (Deploy application), JOB TYPE (Run), LAUNCHED BY (admin), INVENTORY (Development Cloud), PROJECT (App deployment), PLAYBOOK (deploy-application.yml), LIMIT (us-east-2b), and INSTANCE GROUP (tower). Below this is an 'EXTRA VARIABLES' section with tabs for 'YAML' and 'JSON'. The main panel, titled 'Deploy application', shows a real-time stream of the playbook execution. It includes a search bar and a list of tasks with their status and output. The tasks are: 1. PLAY [all] (Successful), 2. TASK [Check required permissions] (Successful), 3. TASK [Fix permissions if needed] (Successful), 4. TASK [Deploy updates] (Successful), and 5. TASK [Restart service] (Successful). The stream ends with a 'PLAY RECAP' summary.

Deploy application

PLAYS 1 TASKS 4 HOSTS 2 ELAPSED 00:00:17

SEARCH

1 PLAY [all] ***** 18:52:29

2

3

4 TASK [Check required permissions] ***** 18:52:29

5 ok: [18.218.56.237]

6 ok: [18.224.239.56]

7

8 TASK [Fix permissions if needed] ***** 18:52:30

9 skipping: [18.218.56.237]

10 skipping: [18.224.239.56]

11

12 TASK [Deploy updates] ***** 18:52:30

13 changed: [18.224.239.56]

14 changed: [18.218.56.237]

15

16 TASK [Restart service] ***** 18:52:33

17 changed: [18.224.239.56]

18 changed: [18.218.56.237]

19

20 PLAY RECAP ***** 18:52:35

21 18.218.56.237 : ok=3 changed=2 unreachable=0 failed=0

22 18.224.239.56 : ok=3 changed=2 unreachable=0 failed=0

23

Real-time status updates

Playbook streams in real-time so you can see success and failures of tasks and jobs.

<https://www.ansible.com/products/tower>

RedHat Ansible Tower

The screenshot displays the RedHat Ansible Tower web interface. The top navigation bar shows the 'TOWER' logo, a user profile for 'admin', and several utility icons. The main content area is divided into two panels. The left panel, titled 'DETAILS', shows the job status as 'Successful' with a green dot. It lists key information: 'STARTED' at 1/23/2019 6:52:19 PM, 'FINISHED' at 1/23/2019 6:52:35 PM, 'JOB TEMPLATE' as 'Deploy application', 'JOB TYPE' as 'Run', 'LAUNCHED BY' as 'admin', 'INVENTORY' as 'Development Cloud', 'PROJECT' as 'App deployment', 'PLAYBOOK' as 'deploy-application.yml', 'LIMIT' as 'us-east-2b', and 'INSTANCE GROUP' as 'tower'. Below this, there are tabs for 'EXTRA VARIABLES' in 'YAML' and 'JSON' formats, with a small table showing one variable. The right panel, titled 'Deploy application', shows a summary of the job execution. It includes a search bar and a table of results. The table has columns for line number, description, and time. The results show a successful execution of the 'Deploy application' job, with a summary at the bottom indicating 'ok=3', 'changed=2', 'unreachable=0', and 'failed=0'.

Line	Description	Time
1	PLAY [all]	18:52:29
2	TASK [Check required permissions]	18:52:29
3	ok: [18.218.56.237]	
4	ok: [18.224.239.56]	
5	TASK [Fix permissions if needed]	18:52:30
6	skipping: [18.218.56.237]	
7	skipping: [18.224.239.56]	
8	TASK [Deploy updates]	18:52:30
9	changed: [18.224.239.56]	
10	changed: [18.218.56.237]	
11	TASK [Restart service]	18:52:33
12	changed: [18.224.239.56]	
13	changed: [18.218.56.237]	
14	PLAY RECAP	18:52:35
15	18.218.56.237 : ok=3 changed=2 unreachable=0 failed=0	
16	18.224.239.56 : ok=3 changed=2 unreachable=0 failed=0	

Activity Stream

See **who** ran **what** job **when**

Complete audit trail of all changes made to Ansible Tower itself - job creation, inventory changes, credential storage, all securely tracked.

All audit and log information can be sent to your external logging and analytics provider

<https://www.ansible.com/products/tower>

CONCLUSION

- Use the right tool for the right job
- Terraform for Infrastructure
- Ansible for Configuration Management
- Leverage Enterprise Solutions

